

Sydney Housing Price Prediction and Decision Support System

SIT720 Machine Learning, Task 8.1D (Distinction)

September 2026

Contents

Code, data and application	2
Part 1: Problem definition and data collection	2
Part 2: Data understanding and feature engineering	4
Part 3: Model development and evaluation	7
Part 4: Investigating prediction failures	10
Part 5: Human judgement, machine learning and LLMs	12
Part 6: Deployment and reflection	14
Reflection	17
GenAI acknowledgement	17
References	18

Code, data and application

GitHub: <https://github.com/BHASKAR-02/sydney-housing-price-prediction>

Live application: <https://<your-app>.streamlit.app>

The repository holds the dataset, the notebook, the FastAPI backend, the Streamlit app and build instructions. `notebooks/sydney_housing_analysis.ipynb` runs top to bottom with no manual steps and reproduces every figure and number below. This report is the summary; the notebook carries the full working.

Part 1: Problem definition and data collection

A real estate agency wants an estimate of a property's sale price from its features, so the target is `sale_price` in AUD and this is supervised regression. The value is speed and consistency: an agent gets a defensible starting number in seconds and identical inputs always give the same answer. It is a decision support tool, not a valuation service, and Parts 4 and 5 show why that distinction matters.

Suburb selection. Three suburbs I would consider buying in, spread across different markets.

Suburb	Postcode	To CBD	Character
Mosman	2088	6.5 km	Harbourside prestige. Large blocks, heritage homes, many water views.
Parramatta	2150	23 km	Sydney's second CBD. Heavy apartment construction, major interchange.
Blacktown	2148	36 km	Outer west, affordable family market. Larger land, lower price per sqm.

The spread matters. Mosman sells on land value and outlook, Blacktown on land size and bedroom count. Parramatta is genuinely split: a new apartment beside the station and an old brick house ten minutes away are two markets sharing a postcode. One global model has to handle all three, and that tension drives the rest of the project.

How the data was collected, and the limitation to read first. Both `realestate.com.au` and `domain.com.au` block automated collection in their terms of use and render sold listings client side, so I could not assemble a few hundred records cleanly. `src/build_dataset.py` generates the dataset instead, from a hedonic pricing structure calibrated against published median sale prices per suburb and property type for 2023 to 2025, with feature distributions matched to sold listings I browsed by hand.

This changes how every accuracy figure here should be read. The data follows a known structure, so the models face an easier problem than real sales would pose. **I would expect the error**

rates below to be noticeably better than this pipeline would achieve on genuinely collected data. The workflow and the conclusions about model behaviour hold; the accuracy numbers are an upper bound. Column names and ranges match what you would scrape, so a real CSV drops in unchanged.

Dataset: 120 properties, 40 per suburb, clearing the 100 total and 30 per suburb minimum. 23 columns covering physical attributes, condition and amenity flags, location, sale context, the agent's blurb as text, and the target. The full schema is in `data/data_dictionary.md`.

Quality, bias and limitations. Missing values are not random: `land_size_sqm` (6 percent) goes missing mostly on apartments where it is meaningless, `year_built` (9 percent) on older stock. I set apartment land to zero explicitly, a real value rather than an absence, and impute the rest by median within suburb and type. Everything here *sold*, so withdrawn listings are invisible and the model will over-predict properties that would struggle to find a buyer. Sales span 24 months during which prices moved, hence the `months_since_start` feature. `suburb`, `distance_to_cbd_km` and `school_catchment_rating` each take three values, one per suburb, so they are the same information three times over, which matters in Part 3. And 120 rows is small for this many features.

Part 2: Data understanding and feature engineering

Price is heavily right skewed, the log is close to normal




Figure 1: Price is heavily right skewed; the log is close to normal

Price skew is 2.03, the usual shape for housing. Logging brings it to 0.61: not perfectly symmetric, but what remains is the handful of Mosman properties above \$5m, which stay unusual on any scale.

Decision: model $\log(\text{sale_price})$ and exponentiate back to dollars. Squared error on raw dollars treats a \$200k miss on a \$5m house the same as on a \$500k unit, but the second is a catastrophe and the first a rounding error; logs make the model care about proportional error, which is how people judge a valuation. Price drivers are also multiplicative (a water view adds a percentage, not a fixed sum), so logging turns them into additive effects a linear model can represent. All metrics are reported back in dollars.




Figure 2: Price by suburb, and by suburb and property type

Mosman's median sits about three times above the other two with a far wider spread. The second panel matters more: within every suburb the type ordering is identical, but the *gap* between types changes by suburb. That interaction suggested a purely additive model would struggle.

The collinearity problem, in numbers. Ranking features by Spearman correlation with price puts `school_catchment_rating` top at **+0.616** and `distance_to_cbd_km` bottom at **-0.616**. Those being exactly equal and opposite is not a coincidence: both take three values, one per suburb, so to

a rank correlation they are one variable with the sign flipped. School rating is not measuring schools, it is measuring which suburb the property is in. This confirms prediction 1 below, and it means an unregularised linear model would be unstable here, splitting one effect across three interchangeable columns. That is a direct argument for the L2 penalty in Ridge. The genuinely independent signals are `internal_area_sqm` (0.530) and `land_size_sqm` (0.499), which vary *within* suburbs.




Figure 3: Price relative to the median of the same suburb and type

Outliers: I kept all of them. None are data entry errors; they are unusual properties that genuinely sold for what they sold for. A model that has never seen an \$8m Mosman house is useless the first time it is asked about one. Dropping real sales improves cross validation scores while making the deployed tool worse. They reappear in Part 4.

My three predictions before engineering anything. (1) **Suburb**, because Sydney is a collection of local markets and no combination of bedrooms bridges a five-fold gap. (2) **Size**, meaning land for houses and internal area for apartments; I say size rather than bedrooms because bedroom count saturates. (3) **Property type**, which changes both the price level and which other features matter.

Nine engineered features: `property_age`, `total_rooms`, `bath_per_bed`, `land_to_internal`, `months_since_start`, `sale_month`, `station_convenience`, `amenity_score`, `description_length`, each justified in the notebook.

Feature	Spearman	Feature	Spearman
<code>land_to_internal</code>	0.355	<code>sale_month</code>	0.109
<code>total_rooms</code>	0.346	<code>bath_per_bed</code>	-0.101
<code>amenity_score</code>	0.279	<code>property_age</code>	+0.069
<code>description_length</code>	0.268	<code>months_since_start</code>	-0.032
<code>station_convenience</code>	-0.129		

Did this match my expectations? Partly, and the mismatches taught me more than the hits. `land_to_internal` and `total_rooms` came top, fitting prediction 2. But two features carry the **opposite sign to intuition**: closer to a station means a *lower* price, and older properties sell for slightly *more*.




Figure 4: Engineered features against price

One explanation covers both. **Mosman has no train line**, being ferry and bus served, and it also holds the oldest stock and the highest prices, while Blacktown and Parramatta sit on rail and are cheaper and newer. These correlations measure “is this Mosman”, with the sign flipped by that accident. The lesson: **across three very different markets, a univariate correlation with price mostly measures how strongly a feature correlates with being in the expensive suburb.** I should not have judged features this way, and would use partial correlations within suburb next time. amenity_score at 0.279 looks strong for the same reason, since Mosman holds the water views and pools.

Part 3: Model development and evaluation

Ridge on log price, the classic hedonic model, chosen for readable coefficients and because the L2 penalty handles the collinearity above; its weakness is that it is additive and cannot represent the suburb-by-type interaction unless I engineer it. **Random Forest**, which gets interactions free and tolerates the outliers I kept, but cannot extrapolate past its training range. **Histogram Gradient Boosting**, usually strongest on small tabular data but easiest to overfit.

My prediction: Random Forest first, Gradient Boosting second, Ridge third, reasoning that the interaction is the biggest structure in the data and that at $n=120$ bagging reduces variance (my problem) while boosting reduces bias (not my problem).

All three sit in a Pipeline so every transformation fits on the training fold only; imputing or fitting TF-IDF before splitting would leak test information. 5-fold rather than 10, since 12 properties per fold would swing the metrics.

Model	MAE	RMSE	R ²	MAPE	Fold-to-fold MAE std	Train R ²	Gap
Ridge Regression	\$281,822	\$456,404	0.920	15.9%	\$58,582	0.975	0.063
Random Forest	\$351,996	\$655,144	0.836	18.9%	\$111,079	0.969	0.123
Gradient Boosting	\$352,600	\$698,258	0.813	17.1%	\$134,170	1.000	0.171




Figure 5: Train against cross validated R squared

Gradient Boosting reaches a training R² of exactly 1.000. It reproduced all 96 training prices perfectly. That is not a model that learned how property is priced, it is a lookup table, and the

0.171 drop is the bill. Ridge has less than half the gap of either tree model, which is what a heavily penalised linear model should look like at this sample size.




Figure 6: Complexity curve for Random Forest depth

The complexity curve shows the same thing directly. At max_depth=2 both curves are low and close: underfitting. As depth rises both climb, then the training curve keeps going toward perfect while the held out curve flattens. Everything past that point is memorisation.




Figure 7: Cross validated predictions against actual

This chart explains the result. All three handle the crowded \$400k to \$2m region well. The difference is entirely at the top. Ridge tracks the diagonal the whole way with mild under prediction past \$5m; both tree models scatter badly, over-predicting some \$2 to \$3m properties by more than a million and under-predicting the most expensive. The cause is structural: a tree predicts the mean of the rows in a leaf, so its output is bounded by the range it has seen. Ridge fits a smooth surface in log space and keeps extrapolating. **With this few expensive properties, extrapolating matters more than modelling interactions.**




Figure 8: Random Forest feature importance

Feature importance confirmed the three predictions but exposed a flaw. The top feature is `txt__mosman`, a TF-IDF term, at 0.21. Because every blurb begins “4 bedroom house in Mosman”, **the text block is re-encoding the suburb I already pass in as a categorical**. Not target leakage, but my TF-IDF block adds far less new information than that number suggests. I would strip suburb and type names before vectorising next time.

Was I right? No, I had it exactly backwards. Ridge won on every metric by about \$70k of MAE. I said I would not claim a winner on a gap smaller than the fold-to-fold spread: Ridge’s MAE standard deviation is \$58,582 and it beats Random Forest by \$70,174, so the gap exceeds its own variability, and Ridge is also the most stable of the three. The result holds.

Three mistakes. I treated the suburb-by-type interaction as the main prize when it is second order and already handed to Ridge through the engineered features. More importantly **I did not think about extrapolation**: the data spans \$405k to \$7.9m with few rows at the top, so much of the error budget falls exactly where trees are structurally incapable. And I chose a log target *because* I argued price drivers are multiplicative, which makes log price near-additive, which is precisely what Ridge fits. I built the transformation that makes the linear model correct, then predicted it would lose.

Recommendation: deploy Ridge. Best on all four metrics, smallest generalisation gap, most stable across folds, it extrapolates, and its coefficients let an agent say “your renovation adds about 7 percent” rather than “600 trees voted”. At twenty suburbs and thousands of sales I would re-run this, since the tree models’ weakness here would shrink.

Part 4: Investigating prediction failures

I rank by percentage error, not dollars: a \$400k miss on an \$8m house is 5 percent and fine; the same miss on a \$600k unit is a disaster. **All five worst errors are over-predictions, and none are expensive properties.**

#	ID	Suburb / type	Actual	Predicted	Error	vs peer median
1	MOS-016	Mosman Apartment	\$900,000	\$1,592,402	+76.9%	0.54x
2	PAR-037	Parramatta Townhouse	\$775,000	\$1,253,851	+61.8%	0.91x
3	BLA-018	Blacktown House	\$720,000	\$1,089,542	+51.3%	0.60x
4	BLA-030	Blacktown Townhouse	\$620,000	\$932,659	+50.4%	0.77x
5	PAR-028	Parramatta Townhouse	\$420,000	\$602,597	+43.5%	0.49x




Figure 9: Error analysis by price level and suburb

In Part 3 I expected the big failures to be under-predicted Mosman waterfronts, because that is where the tree models broke. **That was wrong for the deployed model.** Ridge extrapolates, so it handles the top adequately; its failures are all at the bottom, and every one sold below its own peer median.

Three patterns. These are outliers *within* their peer group: MOS-016 sold at 0.54x what a Mosman apartment normally fetches, so the model priced it correctly as a Mosman apartment and the market disagreed, because whatever made it cheap is in none of my columns. Percentage error is brutal at the cheap end, which is correct, since telling a first home buyer a \$620k townhouse is worth \$930k is far more damaging than a proportional miss at the top. And **the amenity flags actively misled the model**: BLA-018 is marked renovated=1 and has_pool=1, pushing the prediction up, yet sold at 0.60x the Blacktown median, because a binary flag treats a cosmetic refresh and a full rebuild identically.

Three of the five are townhouses, well above their share.

What would have fixed this, in order of value: condition on a 1 to 5 scale rather than binary, since every failure sold below its peer group; exact location, because all of Mosman is one place to this model; floor level and aspect for apartments; vendor motivation; strata levies; and recent nearby comparable sales.

Treat predictions cautiously for cheap properties in expensive suburbs, townhouses, anything outside the 5th to 95th percentile, unusual feature combinations, and anything outside the three suburbs, where the model returns a confident meaningless number.

Are some types inherently harder? Median absolute percentage error:

	Apartment	House	Townhouse
Blacktown	16.4%	11.3%	12.0%
Mosman	16.9%	13.1%	21.7%
Parramatta	13.5%	10.3%	20.4%

Houses are easiest and townhouses hardest, the opposite of what I expected, having assumed apartments would be easiest as near-commodities. Houses win because they are the largest group and their price is dominated by land and internal area, the two things I measure best. Townhouses are hardest because there are only 8 to 10 per suburb *and* they are genuinely ambiguous: land like a house, valued like an apartment, with the balance set by strata arrangements my features cannot express. Apartments sit in the middle because without floor level or aspect, two units with identical bedroom counts can sell 40 percent apart.

This drove a design decision rather than a caveat: the application computes its confidence range **per suburb and property type**, widest for townhouses and Mosman.

Part 5: Human judgement, machine learning and LLMs

I split off the held out set first, fit the model on the rest, and recorded my own estimates and the LLM's **before** seeing actual prices or model output, since anchoring would void the comparison. LLM estimates came from Claude in a separate session with no access to the dataset, one property at a time. My own method was a price per square metre calculation from the Part 2 tables, adjusted for view, condition and station distance.

The estimates live in `data/part5_estimates.csv` rather than being hard coded, so the notebook runs end to end for a marker without a live LLM session, and the ten property IDs are fixed by the `random_state=42` split so they do not move between runs.

Approach	MAE	RMSE	MAPE	Median abs % err	Bias	Within 10%
ML model	\$239,641	\$307,691	18.3%	13.5%	+14.5%	50%
LLM	\$475,500	\$600,206	37.8%	25.9%	+23.3%	20%
Me (human)	\$330,500	\$494,214	22.2%	13.0%	+18.6%	40%




Figure 10: Ten held out properties, three approaches

The model wins the headline metrics, but the detail matters more. My median absolute percentage error is 13.0 against the model's 13.5. On the *typical* property I was marginally better. Its advantage appears in MAE and RMSE, which averages pull around, and in the within 10 percent count. **It does not beat me by being better on the normal house; it beats me by not making the occasional disaster I do.**

All three over-predict by 14 to 23 percent, which for the model is the Part 4 pattern showing through. **The LLM fails in a diagnosable way:** twice as bad on every metric, 2 of 10 within 10 percent. It clearly knows Mosman is expensive and Blacktown is not, so Sydney price levels are in its training data; its failure is pulling every estimate toward the suburb median, so a modest Mosman apartment comes in far too high and a large house too low. It reasons from a remembered median rather than the property in front of it. The dangerous part is the confidence rather than the inaccuracy, since a wrong number delivered with no hedging is worse than a wide range delivered honestly.

Does human judgement still add value? Yes, but not where people usually claim. On ordinary properties I matched the model and lost on the tail, so its real advantage is reliability and scale, not

insight. Humans earn their place in three things it cannot do: **knowing when the model is out of its depth** (I can see a one bedroom Mosman apartment coming and it cannot, at a 77 percent error); **using information not in any column**, since “deceased estate” is a sentence to me and 40 TF-IDF terms to it; and **accountability**, because someone must answer for a wrong valuation and that cannot be a Ridge regression. The sensible design is the model producing the first number and flagging its confidence, with a human reviewing the low confidence cases.

Caveat. Ten properties is tiny and I claim no significance, though the LLM gap is consistent across all six metrics. The 0.5 point median difference between me and the model is well inside noise.

Part 6: Deployment and reflection

Architecture. A **FastAPI** backend loads the Ridge pipeline at startup, validates input with Pydantic, and returns a prediction with a confidence band and warnings. A **Streamlit** front end provides the interface, and can also run the model **in process**, detecting which mode to use. That fallback exists because free Render services sleep after 15 minutes and take about 50 seconds to wake, so a marker opening the link during a cold start would otherwise see an error. It imports the same `predict_one` the API uses, so the paths cannot drift apart.

Two design decisions came straight from the analysis. The 80 percent interval is computed from the residual spread of **that specific suburb and property type**, because Part 4 showed error varies far too much for one global band. And the Part 4 findings are returned as **warnings in the API response** rather than buried in this report.




Figure 11: Entering a property and reading the estimate

This second screenshot is the behaviour I most wanted. Asked about a one bedroom Mosman apartment, the exact profile of MOS-016, the tool returns an estimate **and warns that this is its worst failure case and the number should be read as a ceiling**.

The second tab scores a CSV; only five columns are required and the rest default. The third tab is a model card with the cross validation results, reliable price range, limitations and ethics. Build and deploy instructions are in the repository README.




Figure 12: The model flagging its own worst failure case

Settings

Prediction backend ⓘ

☒ Auto

☐ API

☐ Standalone

Using FastAPI at
<http://localhost:8077>

Deployed model

Ridge Regression

Cross validated error

\$281,822 typical miss

15.9% of sale price

R squared 0.920

Trained on 120 sales across 3 suburbs.

Deploy ⓘ

Sydney Housing Price Estimator

SIT720 task 8.1D. A decision support tool for Mosman, Parramatta and Blacktown. Every estimate comes with a confidence range and its known failure modes, because the model is wrong in predictable ways.

Estimate one property Upload a CSV About the model

Score a spreadsheet of properties

Upload a CSV with one property per row. Missing optional columns are filled with sensible defaults.

suburb, property_type, bedrooms, bathrooms, internal_area_sqm

Download a template CSV

Choose a CSV

Drag and drop file here

Limit 200MB per file • CSV

Browse files

Figure 13: Batch scoring a spreadsheet

Reflection

Hardest part: data collection, and not for technical reasons. The sites do not want to be collected from, fields are inconsistent, and the things that most drive price are written in prose rather than stored as fields. I spent more time deciding what a column should mean than training models.

The simplest model won and I predicted it would come last. Writing that prediction down first is what made the project useful, because it forced me to find where my reasoning broke instead of quietly adjusting my expectations to fit the output. **Complexity does not buy accuracy when data is the constraint:** Gradient Boosting fit its training data perfectly and came last. At 120 rows the binding constraint is information, not capacity, so another week would go on collecting rows and columns, not on hyperparameters. **Feature engineering needs checking, not just doing:** two of my nine features did not work as intended, and I only found the TF-IDF suburb leak by inspecting importances.

Performance versus deployability. I got lucky, since the most accurate model was also the most interpretable. Had Gradient Boosting won narrowly I would still have shipped Ridge, because a valuation an agent cannot explain is worth less than one slightly less accurate and defensible.

Ethics. Two concerns. *Feedback loops:* if agents set asking prices from the tool, those influence sale prices, and those sales become next year's training data, the model ends up validating its own past predictions and any bias gets baked in rather than corrected. *Proxy features:* `school_catchment_rating` and `distance_to_cbd_km` correlate with the income and demographic composition of an area, so a model learning "this postcode sells for less" is partly learning a socioeconomic pattern, and using it to *set* prices rather than describe them entrenches it. Including no explicitly demographic feature does not make the model neutral, because the location features carry that information anyway.

A fairness problem is also visible in my own error analysis: the model is least accurate at both ends of the price range, so it serves first home buyers and prestige vendors worst, and of those two the group least able to absorb a bad valuation is the first. **An MAE of \$281,822 sounds acceptable until you notice it is not evenly distributed across who gets harmed by it.**

With more resources, in order: several thousand sales rather than 120; latitude and longitude so the model learns position within a suburb; a comparable sales feature; real quality grading; stripped suburb names in the text block; and a time series component modelling the market index separately from the property.

The limitation that matters most. Because the data was generated rather than collected, the models faced an easier problem than production would pose. That a linear model reaches R^2 of 0.92 is itself a warning sign, since the generator was close to multiplicative and Ridge on a log target is close to the right functional form almost by construction. Real accuracy would be lower and the gaps between models narrower. The workflow, the code and the conclusions about overfitting and extrapolation hold. The accuracy numbers are an upper bound.

GenAI acknowledgement

I used Claude (Anthropic) for planning the notebook structure, suggesting engineered features to try, debugging a scikit-learn `ColumnTransformer` error where my text column was passed as a `DataFrame` instead of a `Series`, and proofreading this report. I also used it as the LLM valuation system in Part 5, which is what that comparison is about.

All modelling decisions, the suburb selection, the feature engineering rationale, the interpretation of results and the conclusions are my own. I checked every suggestion against the actual output before keeping it, and rejected several, including an early recommendation to drop the outlier properties, which I argue against in Part 2.

References

Australian Bureau of Statistics 2024, *Residential Property Price Indexes: Eight Capital Cities*, ABS, Canberra.

Breiman, L 2001, 'Random forests', *Machine Learning*, vol. 45, no. 1, pp. 5 to 32.

Hastie, T, Tibshirani, R & Friedman, J 2009, *The Elements of Statistical Learning*, 2nd edn, Springer, New York.

NSW Valuer General 2024, *Property Sales Information*, NSW Government, Sydney.

Pedregosa, F et al. 2011, 'Scikit-learn: Machine learning in Python', *Journal of Machine Learning Research*, vol. 12, pp. 2825 to 2830.

Rosen, S 1974, 'Hedonic prices and implicit markets', *Journal of Political Economy*, vol. 82, no. 1, pp. 34 to 55.